

MERIDIAN TRUST & ASSET SERVICING, N.A.

Personal Fiduciary Administration · First Line of Defense

OPS-NOTE-2147-01

Self-Identified Open Items

Discretionary Distribution & Fiduciary Review Intelligence Assist (DDFR-IA)
v1.4

Document number	OPS-NOTE-2147-01
Version	0.3
Classification	Internal — Restricted
Document owner	Head of Personal Fiduciary Administration & Wealth Advisory
Effective	Current cycle
Review cycle	On request
Approved by	Not approved — working document

Contents

Document Control	1
1. Purpose	1
2. Open Items	1
2.1 No structured accuracy testing has been performed	1
2.2 No adversarial or prompt-injection testing	2
2.3 Data protection impact assessment not completed	2
2.4 No bias or differential-treatment evaluation	
2.5 Retrieval quality is unmeasured	2
2.6 Human review friction is not measured	2
2.7 Third-party terms not confirmed for all data classes	
2.8 No production monitoring	3
2.9 Records of algorithmic contribution are not retained	3
2.10 Registration was retrospective	3
3. Context the team wishes to record	3

Document Control

Document number	OPS-NOTE-2147-01
Document type	Working note — first line
Status	Draft. Not reviewed by Model Risk Management.
Related system	UC-2147 / AI-REG-2026-0148 — DDFR-IA v1.4
Prepared by	Director, Wealth Platform Engineering
Reviewed by	—
Approved by	—

Note on the status of this document. This is a working note maintained by the development team. It is not a validation report, it has not been independently reviewed, and it is not evidence of control effectiveness. It is submitted with the CAIMI registration for UC-2147 because the AISGC requested disclosure of known issues, and it is the only artifact currently available.

1. Purpose

1.1 This note records issues the development team is aware of in DDFR-IA v1.4, which has been operating in limited production across three trust administration teams since the prior quarter.

1.2 It is not a control assessment. Where an item below describes a missing control, the absence is stated; no compensating control is claimed unless explicitly described.

2. Open Items

2.1 No structured accuracy testing has been performed

We have not measured how often the system asserts something the governing instrument does not support. Trust officers have reported individual instances informally — most memorably a memorandum asserting an unconstrained power to invade principal for a business venture where the instrument permitted invasion only for medical expenses — but these are anecdotes collected in passing, not a measurement.

No golden dataset exists. No threshold has been set. We do not currently have a way to say whether the rate is improving or worsening between releases.

2.2 No adversarial or prompt-injection testing

Not performed. The system ingests beneficiary-submitted correspondence and supporting documentation, which is untrusted content from an external party, and we have not tested what happens when that content contains instruction-like text.

2.3 Data protection impact assessment not completed

A DPIA was scoped in the prior quarter and not completed. The system processes beneficiary names and Social Security numbers, family relationship records, personal financial statements, and — for health and education distributions — **medical and educational records**. Source documents are submitted to the model as scanned originals with no pseudonymisation.

2.4 No bias or differential-treatment evaluation

Not performed, in any form.

We understand from STD-AI-002 Control 4 that for a system operating in a fiduciary capacity this requires a population-level measurement of allocation outcomes across beneficiary classes, and that reviewing individual memoranda does not substitute for it. We have not done this and we do not currently have the instrumentation to do it — the system does not retain a structured record of its income-versus-principal allocation recommendations in a form that could be analysed across a population.

One trust officer raised a concern during the pilot that the system "seems to reach for income first." We have no data with which to confirm or refute that.

2.5 Retrieval quality is unmeasured

Context precision and recall have not been measured. There is no golden dataset of instrument-and-question pairs with adjudicated correct passages.

We are aware of a specific structural concern. Trust indentures in the affected population routinely run past 200 pages, with amendments and codicils appended, and our chunking strategy is a fixed sliding window with overlap. We have observed cases where a grant of discretionary authority and the restriction qualifying it were separated across chunks, and only the grant was retrieved. We have not quantified how often this occurs.

2.6 Human review friction is not measured

Each memorandum is placed before the Fiduciary Review Committee. We do not capture modification rate, rejection rate, or time in review, so we cannot evidence that the review is doing work. Committee minutes record outcomes, not the extent of change made to the draft.

Under current volumes each trust officer receives 40–60 draft memoranda per cycle.

2.7 Third-party terms not confirmed for all data classes

The enterprise no-training confirmation with the foundation-model provider is in place. It does not specifically name **health information** or **records concerning minor beneficiaries**, both of which this system processes. We have raised this with Third-Party Risk Management and have not received a response.

2.8 No production monitoring

None. There is no dashboard, no indicator, no threshold, and no alert. Issues surface when a trust officer reports one.

2.9 Records of algorithmic contribution are not retained

The final memorandum is archived to the core trust accounting system. The prompt, the retrieved indenture passages, the model version, and the generation parameters are **not retained**. If asked to demonstrate how a particular recommendation was reached, we could not reconstruct it.

We understand this is the subject of Internal Audit finding IA-2026-031 and of STD-AI-002 § 12.3.

2.10 Registration was retrospective

The system entered limited production before it was registered in CAIMI. Registration was completed this cycle at the direction of the AISGC.

3. Context the team wishes to record

3.1 The pilot was initiated to meet a service-level commitment to the Personal Trust division following an increase in distribution request volume. The team understood at the time that it was a pilot of a drafting aid.

3.2 The team does not dispute any item above and is not asking for them to be waived. Formal testing and documentation are planned but not yet complete, and no date has been committed.

3.3 The team requests guidance on sequencing, given that several items require instrumentation that does not currently exist in the application.

SYNTHETIC DEMONSTRATION DOCUMENT. Meridian Trust & Asset Servicing, N.A. is a fictional institution. This note and the system it describes were authored for a SAS software demonstration.